

# Parallel Fast N-Body

HPCSE II - Spring 2015

# Last time

- N-body solvers for long range interactions
  - Naive:  $O(N^2)$
  - Idea: treat cluster of far away particles as one
    - Barnes-Hut:  $O(N \log N)$
    - Fast Multipole Method:  $O(N)$

# The Barnes-Hut algorithm

- Asymptotic complexity:  $O(N \log N)$
- BUT much higher implementation complexity

tree building

multipole expansions

tree traversal

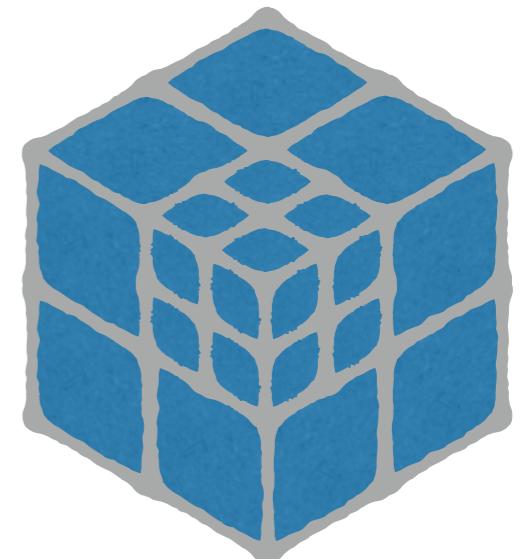
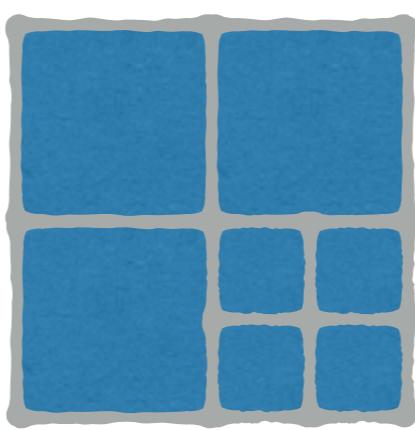
force summation

VS

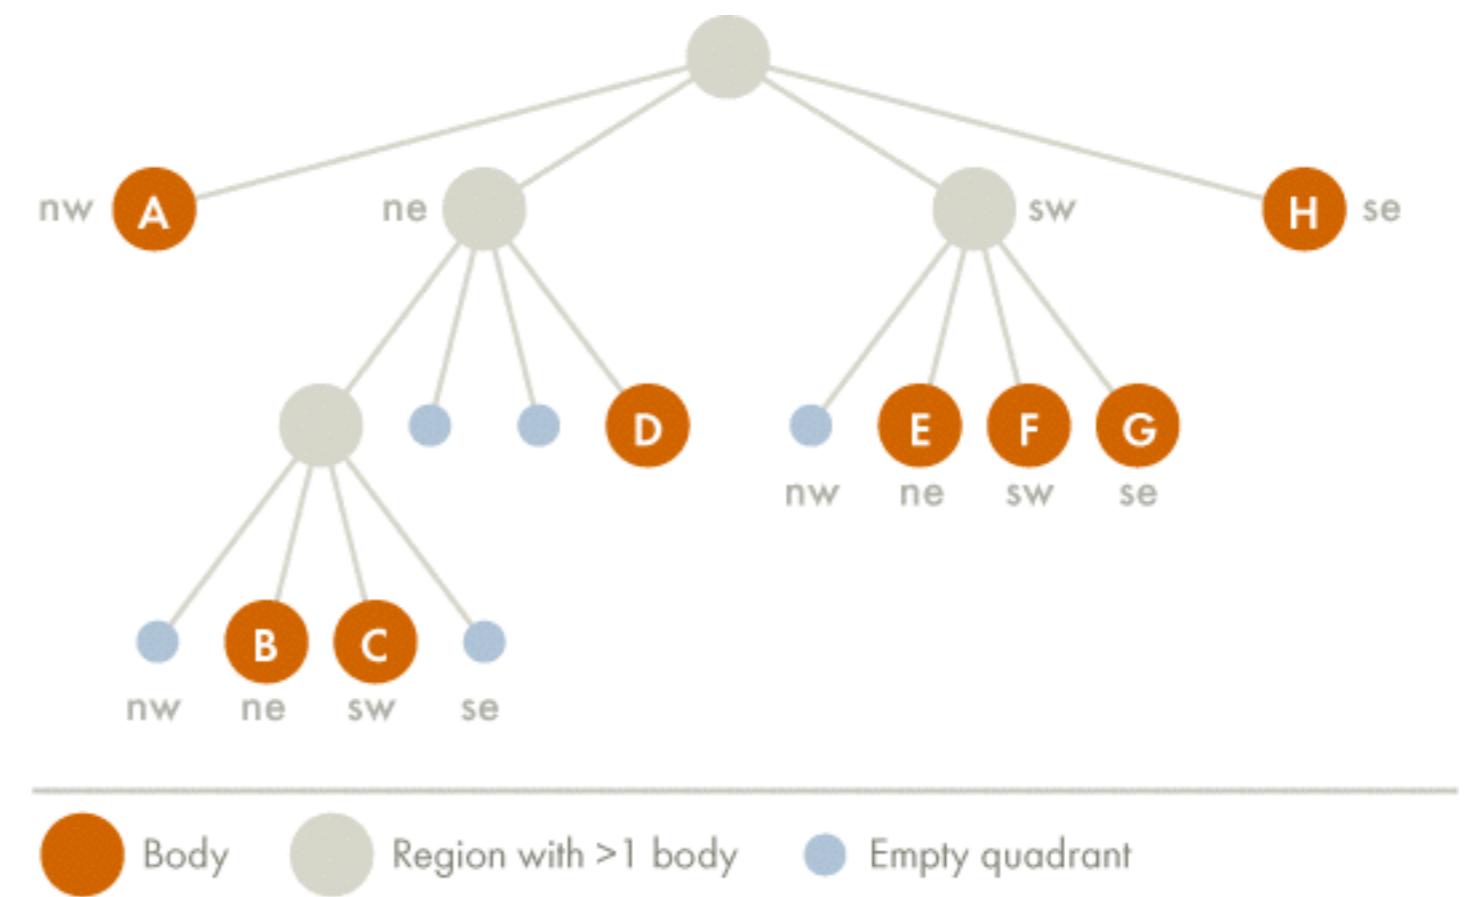
force summation

# Data structure

- Requirements:
  - Hierarchical spatial decomposition
  - Simple to build and traverse
- Extension of the binary tree in multiple dimensions:
  - Quadtree (2D)
  - Octree (3D)



# Quadtree



Body



Region with >1 body



Empty quadrant

# Building the quadtree

- Top-down approach
- Particles are inserted in an initially empty tree
- 3 types of nodes:
  - internal - contain multiple leaves
  - leaf - contains a single particle
  - empty

# Building the quadtree

for each particle

    current\_node = root\_node

    while (true)

        if (current\_node != empty && !internal)

            insert particle, exit loop

        else if (current\_node == internal)

            put into corresponding quadrant of current\_node

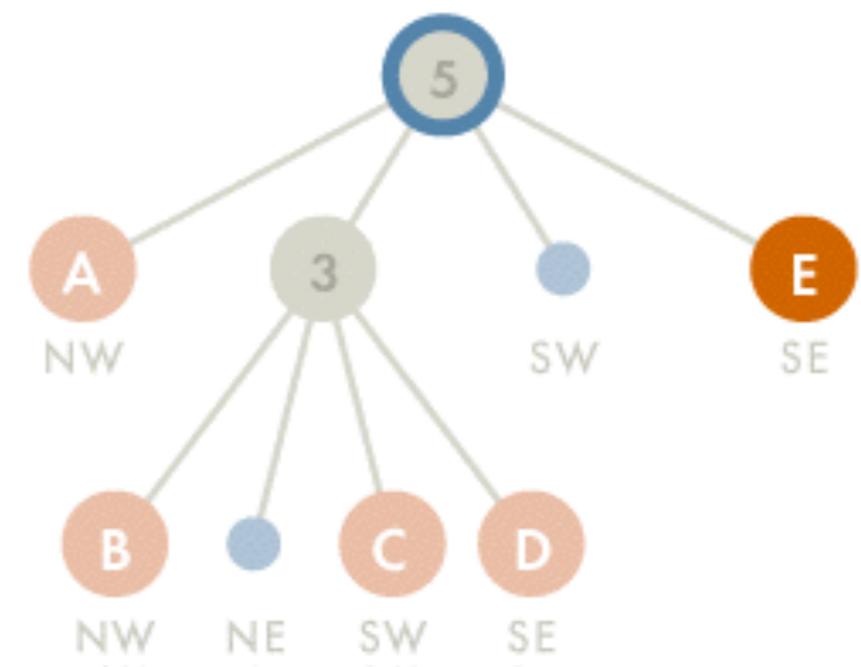
    else

        split current\_node into its children (one for each quadrant)

        put particles into their quadrants nodes

# Building the quadtree

Add body E  
fits in se corner of root node



# Multipole expansions

- Multipole expansions computations
  - Option 1: while building the tree  
(more difficult to parallelize)
  - Option 2: after building the tree  
(bottom-up, “reduction” operation)

# Traversing the tree

for each particle

```
current_node = root_node
```

```
if (current_node==leaf && particle contained not self)
```

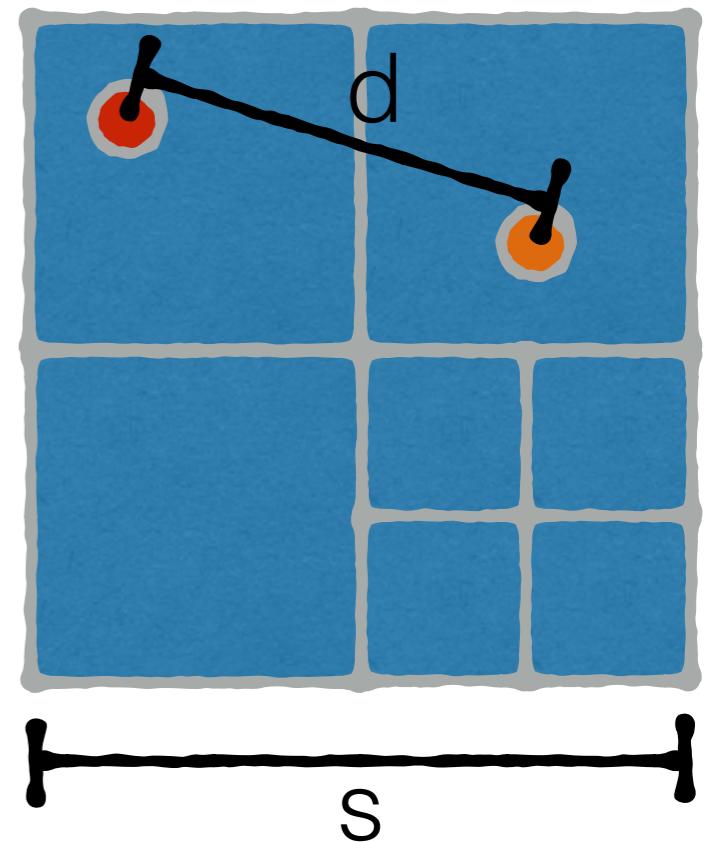
compute force, add to self

```
else if (s/d<theta)
```

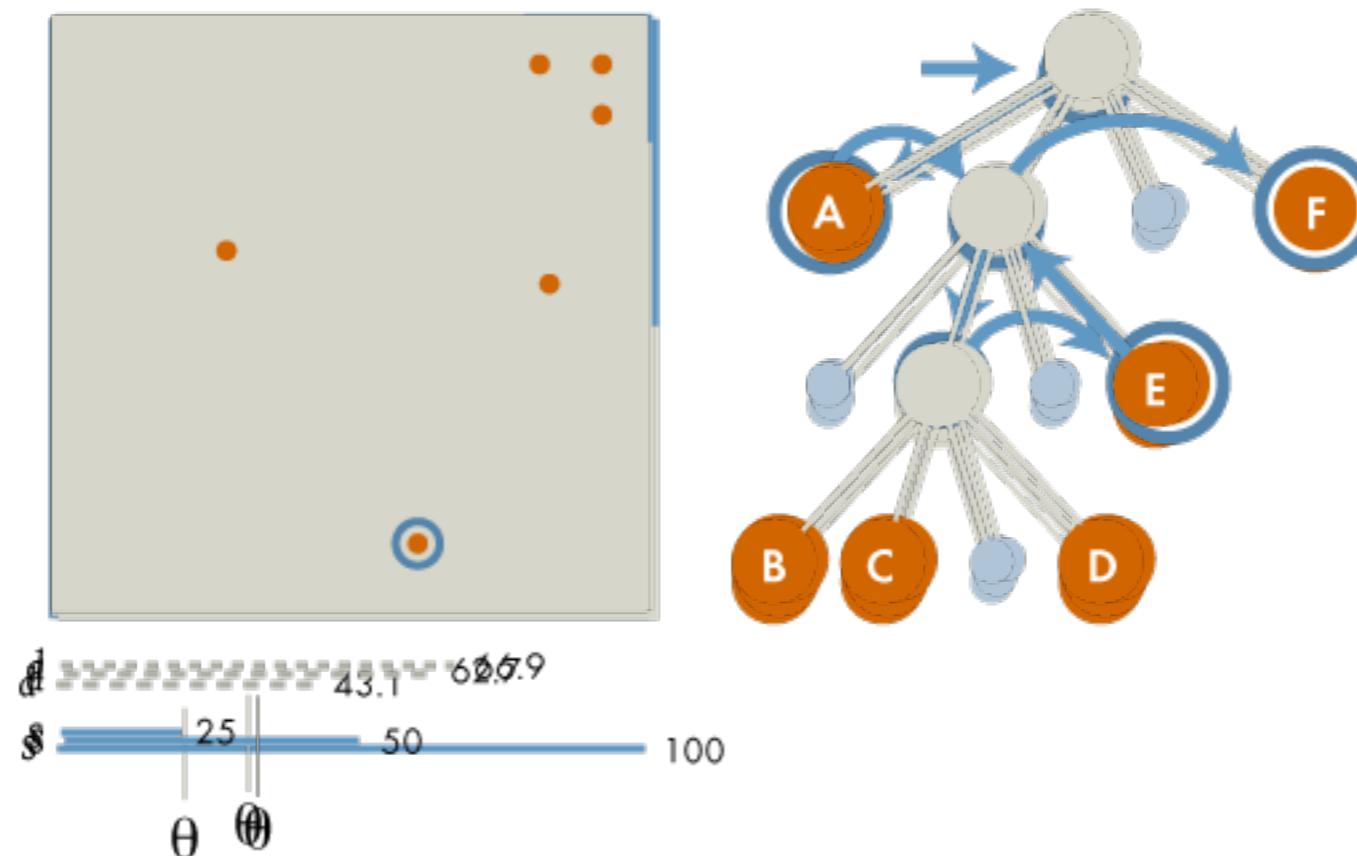
compute force with internal node

```
else
```

follow this procedure on each child



# Traversing the tree

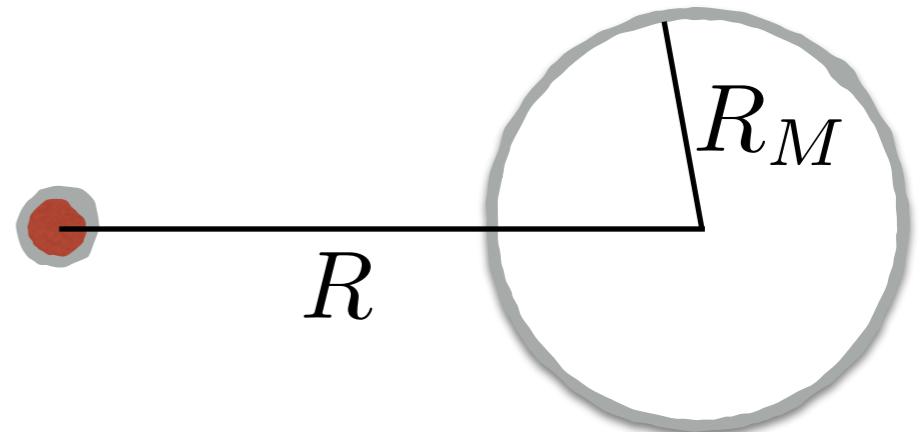


# Fast multipole method

- Particle-box interactions
- Box-box interactions
- Convergence rate proportional to

$$\left(\frac{R_M}{R}\right)^{P+1}$$

P multipole moments expansions



# Box-box interactions

- Interactions computed at the coarsest level

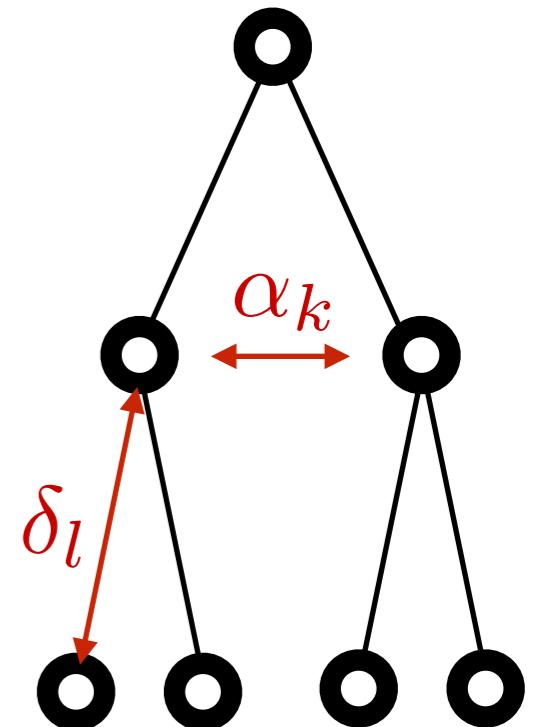
$$\delta_l = \frac{(-1)^{l+1}}{(Z_G - Z_M)^l} \sum_{k=0}^P \binom{l+k-1}{k} \frac{\alpha_k}{(Z_G - Z_M)^k}$$

$$l = 1, \dots, P$$

- Velocities computed at finest level

$$\delta_l^{\text{children}} = \sum_{k=l}^P \binom{k-1}{k-l} \delta_k^{\text{parent}} (Z_M^{\text{children}} - Z_M^{\text{parent}})^{k-l}$$

$$V(Z) = \frac{i}{2\pi} \sum_{l=1}^P \delta_l (Z - Z_G)^{l-1}$$

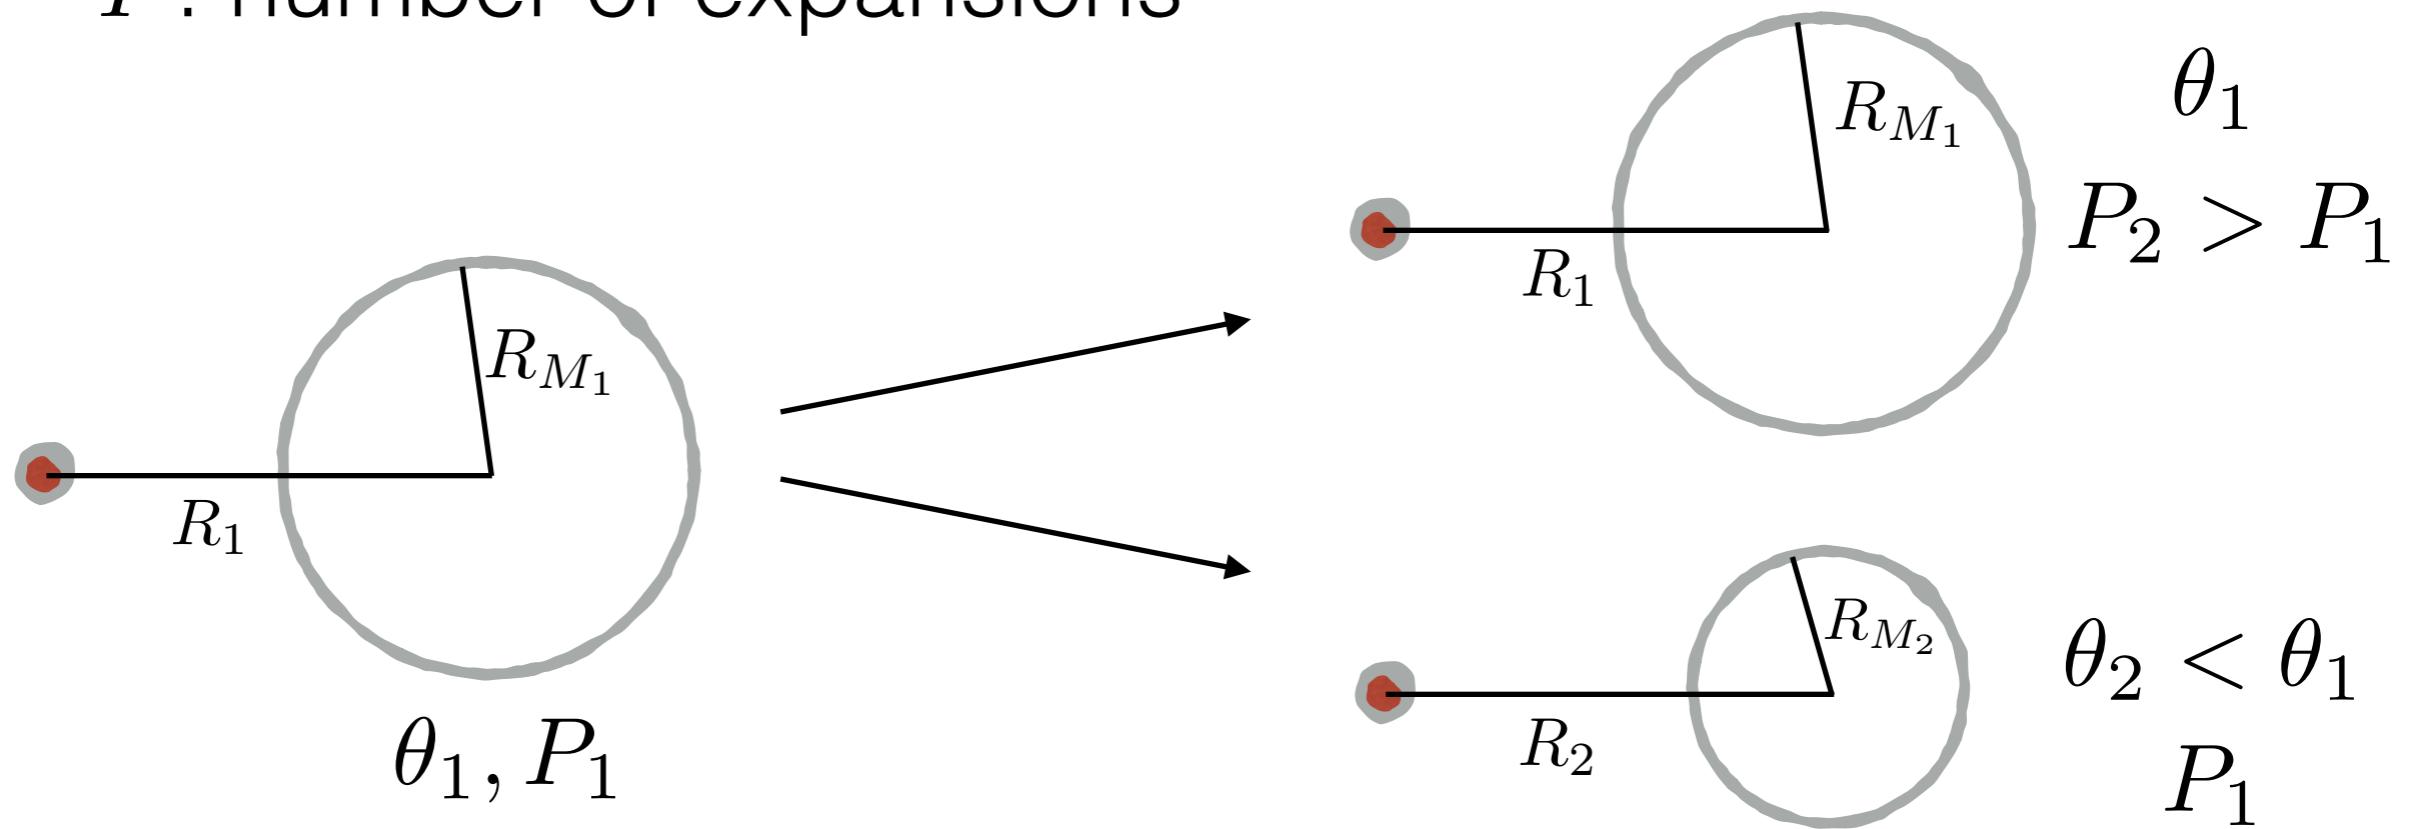


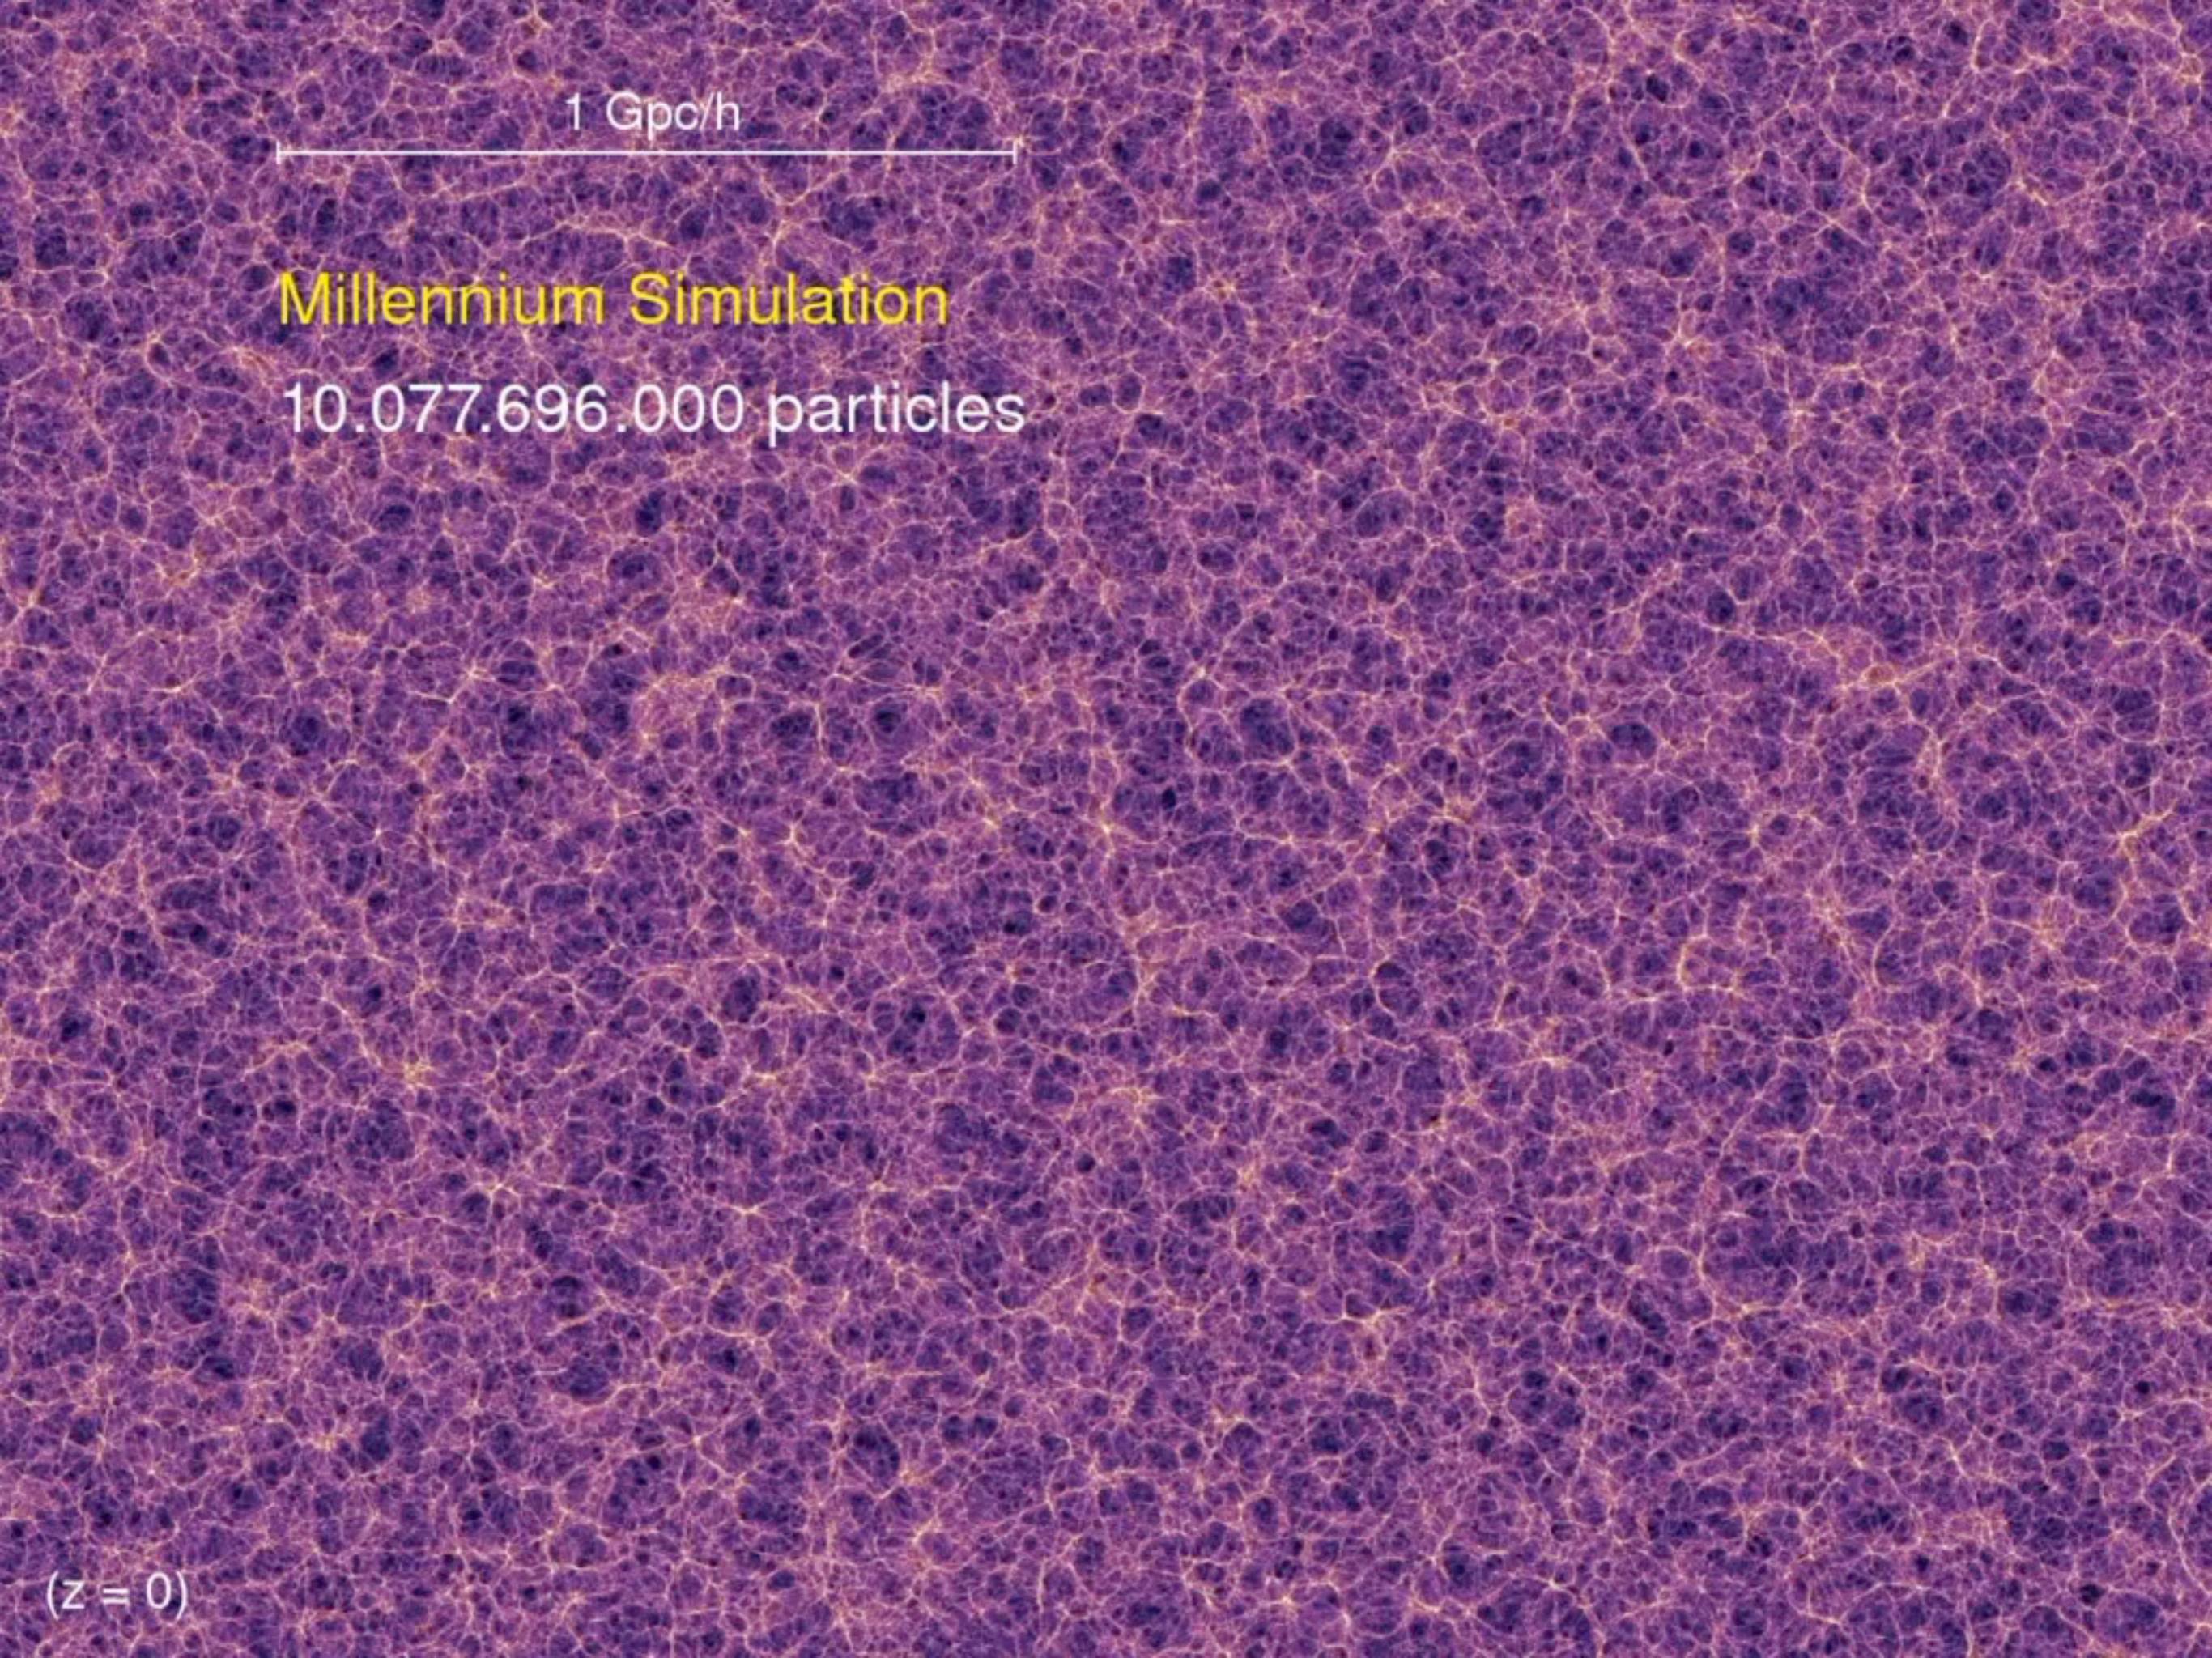
# Accuracy considerations

- Multiple parameters affect accuracy

- $\theta = \frac{R_M}{R}$  : ratio of box size to distance

- $P$ : number of expansions





1 Gpc/h

Millennium Simulation

10.077.696.000 particles

(z = 0)

# Parallelization of BH

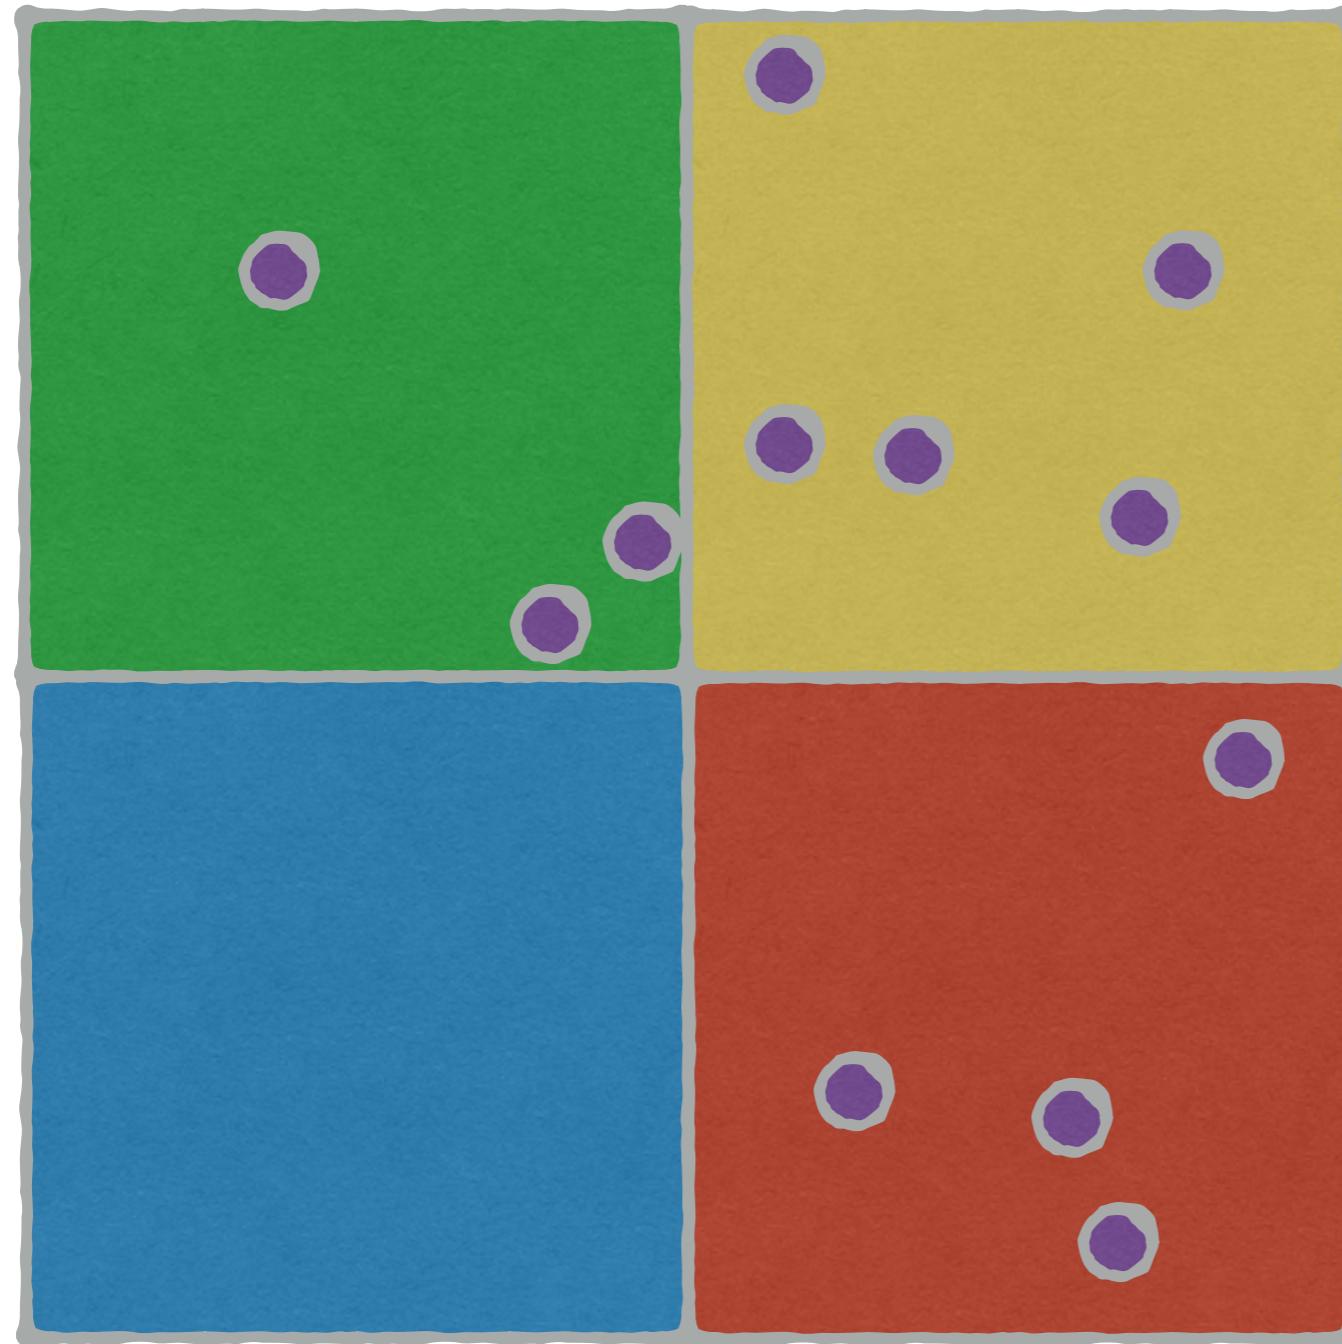
- Material based on Gibbon et al, Parallel Tree Codes ([http://juser.fz-juelich.de/record/16155/files/IAS\\_Series\\_06.pdf](http://juser.fz-juelich.de/record/16155/files/IAS_Series_06.pdf))
- Irregular code
  - How to distribute and define the work?
  - How to minimize communication in MPI?

# Parallel algorithm

for t in time

- (i) perform domain decomposition
- (ii) build tree
- (iii) traverse tree and build interaction list
- (iv) perform force summation

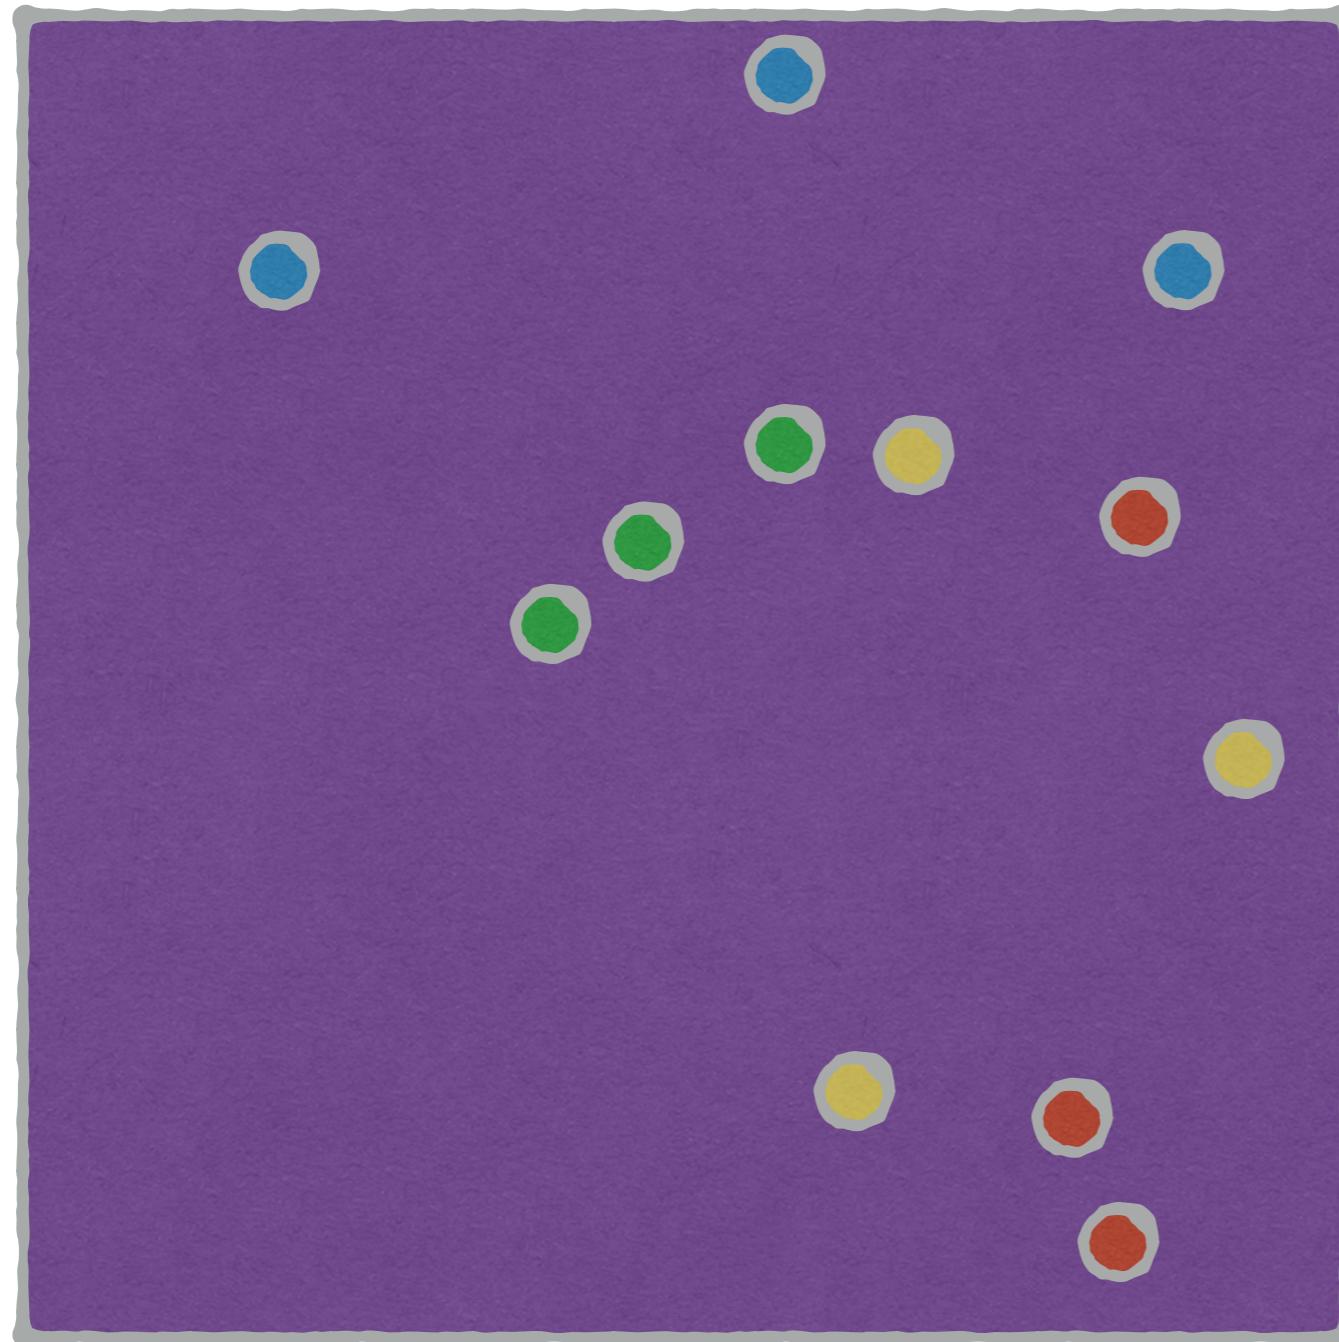
# Domain decomposition



✓ Simple

✗ Highly unbalanced

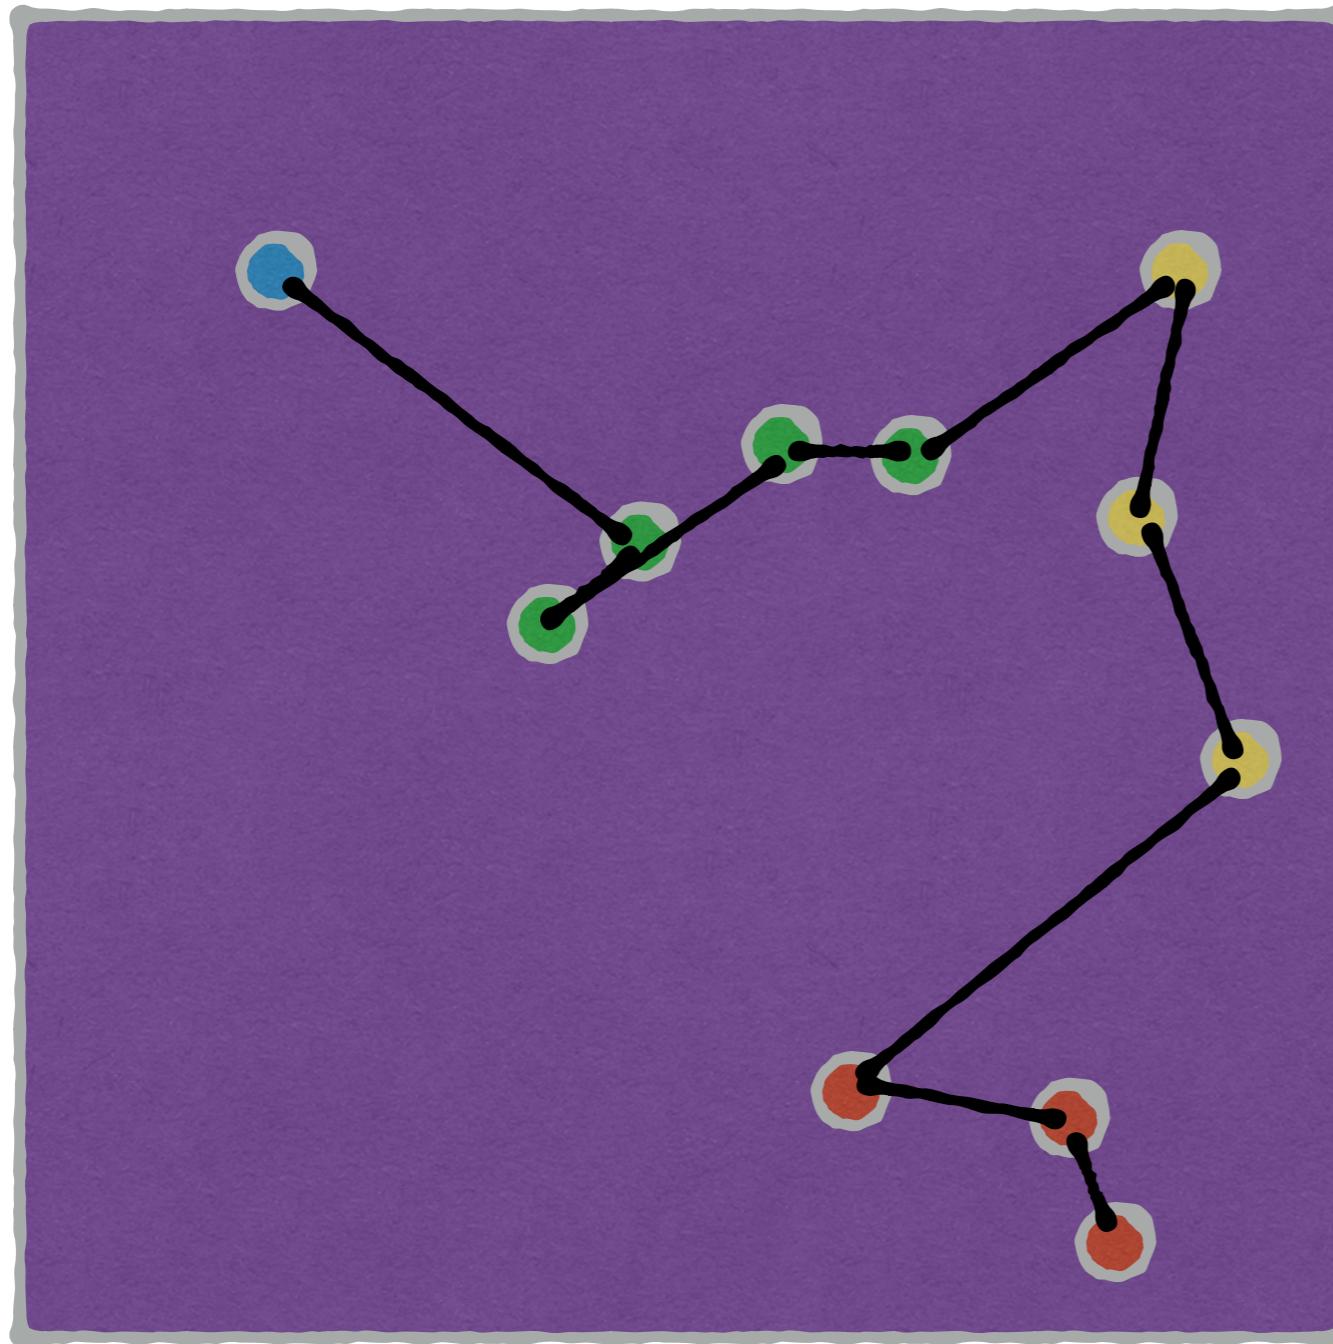
# Domain decomposition



- ✓ Simple
- ✗ Unbalanced  
(different number of force summations)
- ✗ Bad locality

vision - what to treat as multipole expansion and what to treat as direct summation

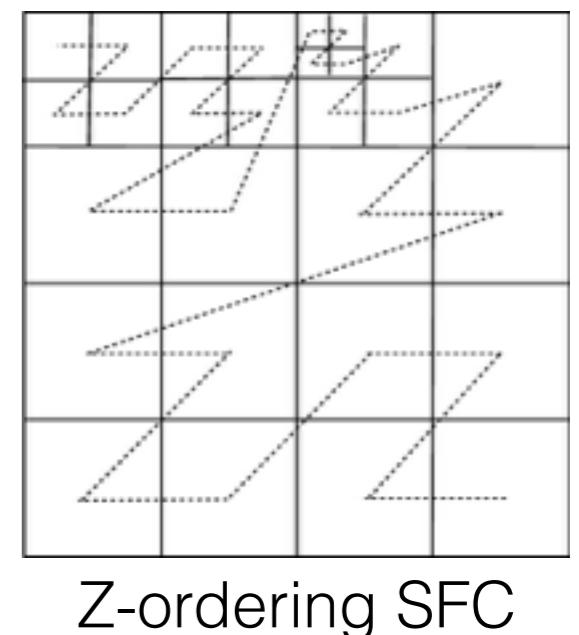
# Domain decomposition



- ✓ Good locality with space-filling curves
- ✓ Load balancing
- ✗ Requires sorting

# Space-filling curves

- SFC show better locality
- Split particles along a 1D curve
- Morton indexing (Z-ordering):  
computation by interleaving bits of coordinates
  - (i) Compute particle index
  - (ii) Sort particles
- Other options: Hilbert, Peano



# Load balancing

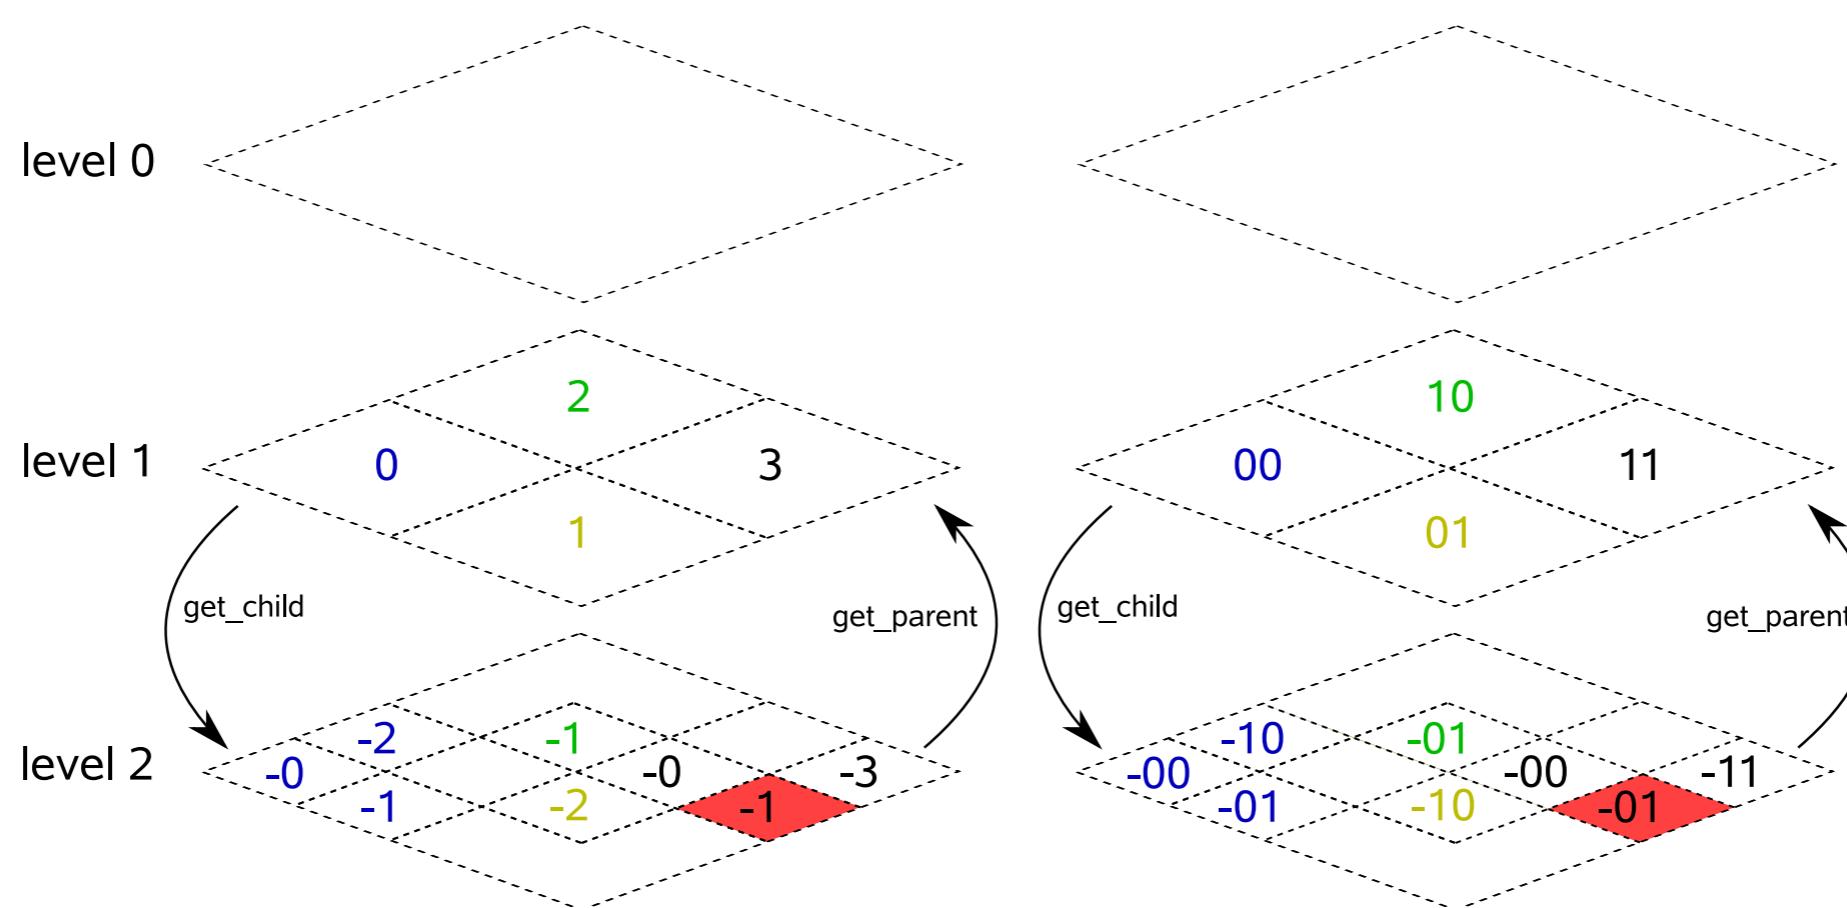
- Use information of previous time steps
  - #interactions computed for a given particle?
  - Rationale: the number doesn't change much

# Building the quadtree

- Each parallel unit builds its part of the tree
- Needed: “ghosts” at extremes of 1D curve
- Z-ordering encodes the quadtree
  - position in the tree must be extracted
  - access is  $O(1)$  instead of  $O(\log N)$

# Building the quadtree

- Parents can be obtained by right bit shifts
- Children are obtained by left bit shifts



picture from Gibbon et al, Parallel Tree Codes

# Tree traversal

- Each parallel element creates an interaction list
- On distributed memory architectures:  
interaction list reduces communication messages

# Force summation

- Given interaction lists, update particles
- Separate step:
  - keeps physics separated
  - some load balancing can be performed